

We'll need to develop a model that can analyze product images and extract specific entity values (like weight, dimensions, etc.) from them.

1. Data Preprocessing:

- Download the images using the provided `download_images` function from `src/utls.py`.
- Organize the training data, matching images with their corresponding labels from `train.csv`.

2. Model Development:

- We'll likely need to use a combination of computer vision and natural language processing techniques.
- A deep learning model, possibly based on a pre-trained convolutional neural network (CNN) architecture like ResNet or EfficientNet, could be used for image feature extraction.
- We might need to incorporate optical character recognition (OCR) to extract text from the images.
- A custom head could be added to the model to predict the entity values based on the extracted features and text.

3. Training:

- Split the training data into training and validation sets.
- Train the model using the labeled data from `train.csv`.
- Use techniques like data augmentation to improve model generalization.

4. Prediction:

- Apply the trained model to the test images from `test.csv`.
- Post-process the model outputs to ensure they match the required format (e.g., "x unit" where x is a float and unit is from the allowed list).

5. Output Generation:

- Create the `test_out.csv` file with the required columns: 'index' and 'prediction'.
- Ensure all test samples have a prediction (use an empty string "" if no value is found).

6. Validation:

- Use the provided sanity.py script to check that the output file passes all formatting checks.

Key Considerations:

- Ensure predictions use only the allowed units specified in constants.py.
- Pay attention to the output format requirements, especially for numerical values and units.
- Handle cases where no value is found in the image by returning an empty string.
- Make sure to generate predictions for all indices in the test set.

Next Steps:

1. Set up the development environment with necessary libraries (e.g., PyTorch or TensorFlow for deep learning, OpenCV for image processing).
2. Implement the data loading and preprocessing pipeline.
3. Design and implement the model architecture.
4. Develop the training loop and evaluation metrics.
5. Create the prediction and output generation scripts.
6. Test the entire pipeline and validate the output using the sanity checker.